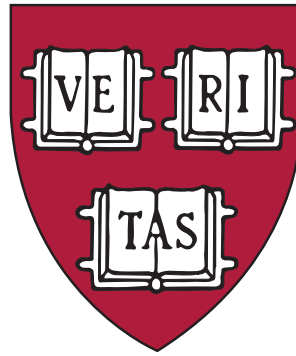


HARVARD UNIVERSITY



Information Technology

Integer - An Integrated Management Environment Design Review - v1.0

June 25, 2014

Topics

- Introduction - Jim Waldo.
- Problem.
- Architecture.
 - Implementation technologies.
 - Logical design.
 - APIs/Interfaces.
 - First release services.
 - Object hierarchies.
- Deployment.
- First release - implementation and operation.
- Contact us and information.

Problem

What's the Problem?

- We don't have a service based view of our systems and services.
- Information gaps cause downtime.
- Non-integrated management systems from multiple sources is not cost-effective
- New environments like AWS:
 - Add complexity.
 - Proprietary methods.

Service Outage Analysis

- No service configuration software.
- Installation and upgrade procedures require significant human interaction, increase error opportunities.
- Effective root cause/systems failure analyses not routinely performed.
- Unknown service dependencies.
- Configuration errors.

Integer's Scope

- Covers essential areas of management with a whole-service view:
 - Configuration:
 - Automated and reliable configuration of systems and services - including deployment of systems & images.
 - Performance:
 - Issues related to latency and capacity.
 - User level to fine-grained per-system details.
 - Fault - detecting and repairing failures.
 - Security - Service wide perspective security controls.
 - Accounting - quantity of work done on behalf of a service.
- Key elements of the environment:
 - Users.
 - Technical Components:
 - Servers.
 - Network elements like routers, firewalls, load balancers, DNS system and other physical and virtual network elements.
 - Software from the virtualization layer to high-level web services.

Integer's Primary Differentiator

- Works with network devices, servers and software as they exist.
- Accepts multiple monitoring and control protocols per system.
- Does not assume managed systems adopt CIM, WebM or any other model or current fad.
- Creates a common representation from these disparate protocols:
 - Enables a services view.
 - Allows cross vendor and technology configuration.

Why Start with Discovery

- Require up to date information on systems/services.
- Time to upgrade existing discovery system.
- No discovery currently provides linkage between network layers.
- Enough of framework to validate architecture and deliver user value.
- Allows early feedback.

Architecture

Architectural and Implementation Focus

Integer is designed to view, manage, and understand the relationships of all the elements of a business service as a whole.

System Overview

- Web front end.
- Services with APIs that hide DB structures.
- Coded in Java.
- A 'core' installation that can be clustered/distributed:
 - Reliability.
 - Performance.
- Multiple distributed server installations.
 - To reduce network impact.
 - Reduce network configuration complexity.

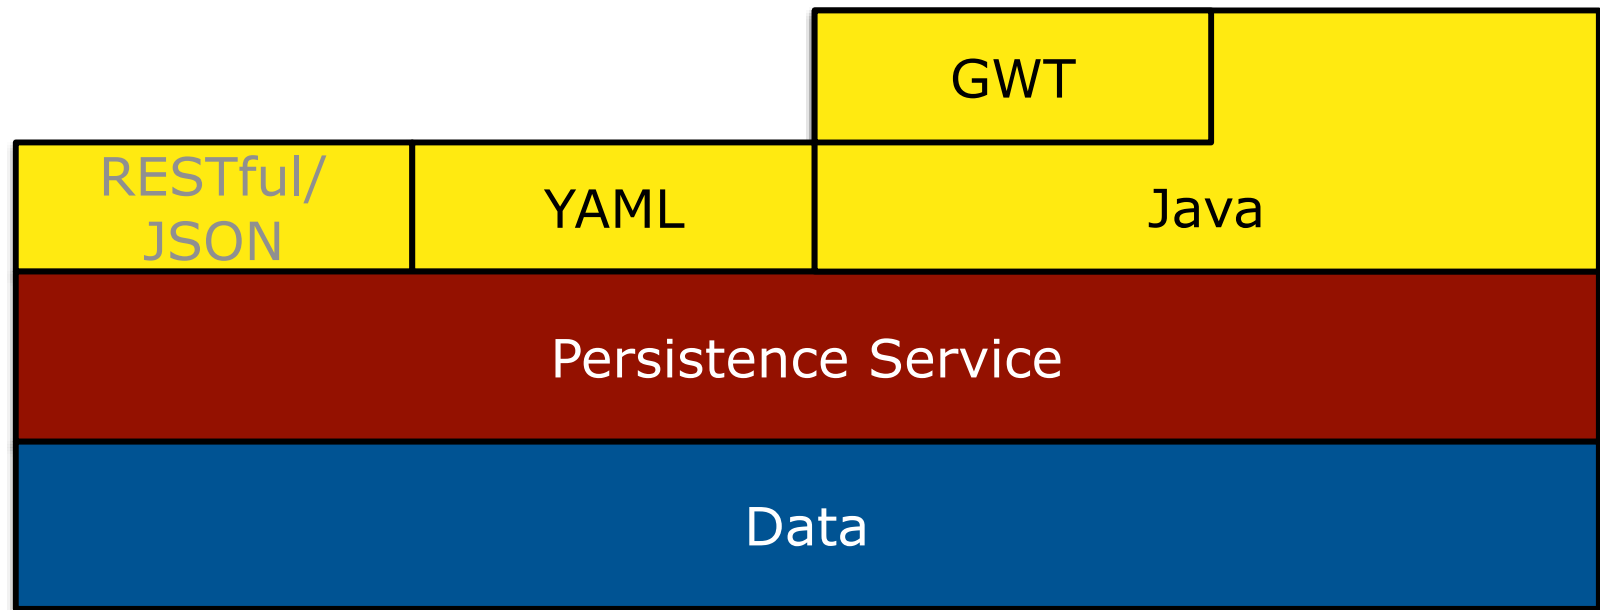
Open Source Technologies/Components

Database	MariaDB
NoSQL	Cassandra - not used in initial discovery release
Middleware	WildFly (a.k.a. JBOSS)
UI	GWT - Google Web Toolkit
Scheduling/Time	Quartz
SNMP	SNMP4j
MIB Compiler/ Loader	Mibble
Logging	SLF4J - Simple Logging Facade for Java

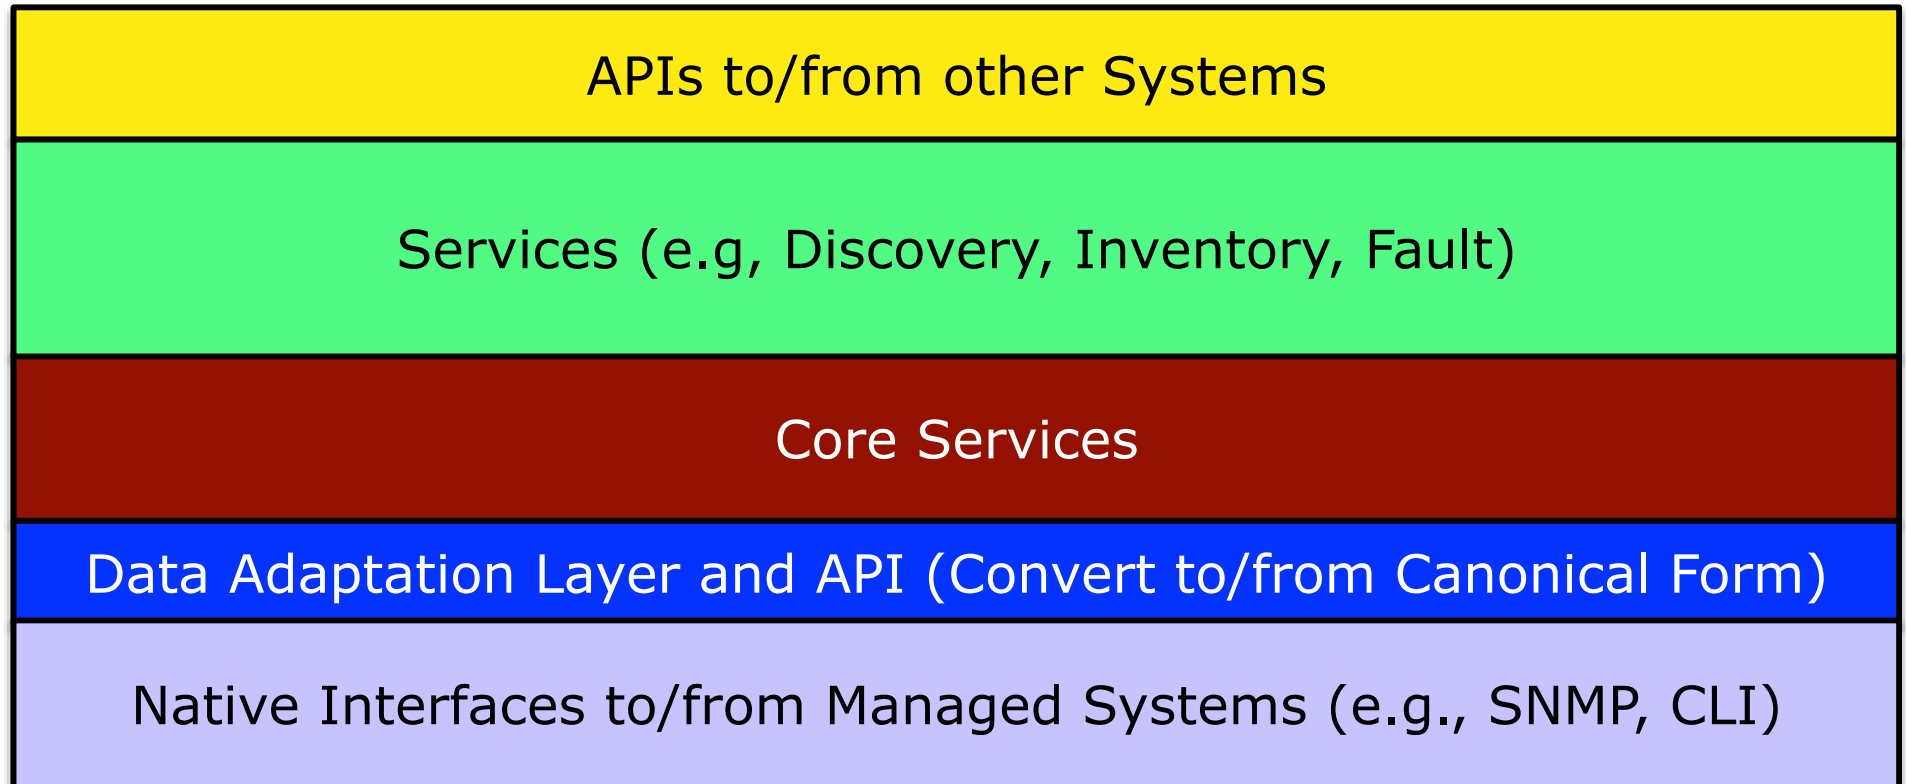
Components We Are Developing

- Persistence service that accesses data in the DB(s):
 - DAO's.
 - APIs to our objects.
- System distribution built on WildFly.
- Interfaces to native interfaces used by devices, e.g., discovery service uses open source SNMP code.
- Our business logic/services, major elements:
 - Discovery
 - Inventory
 - Reporting
- UI built out of GWT with some custom JS.

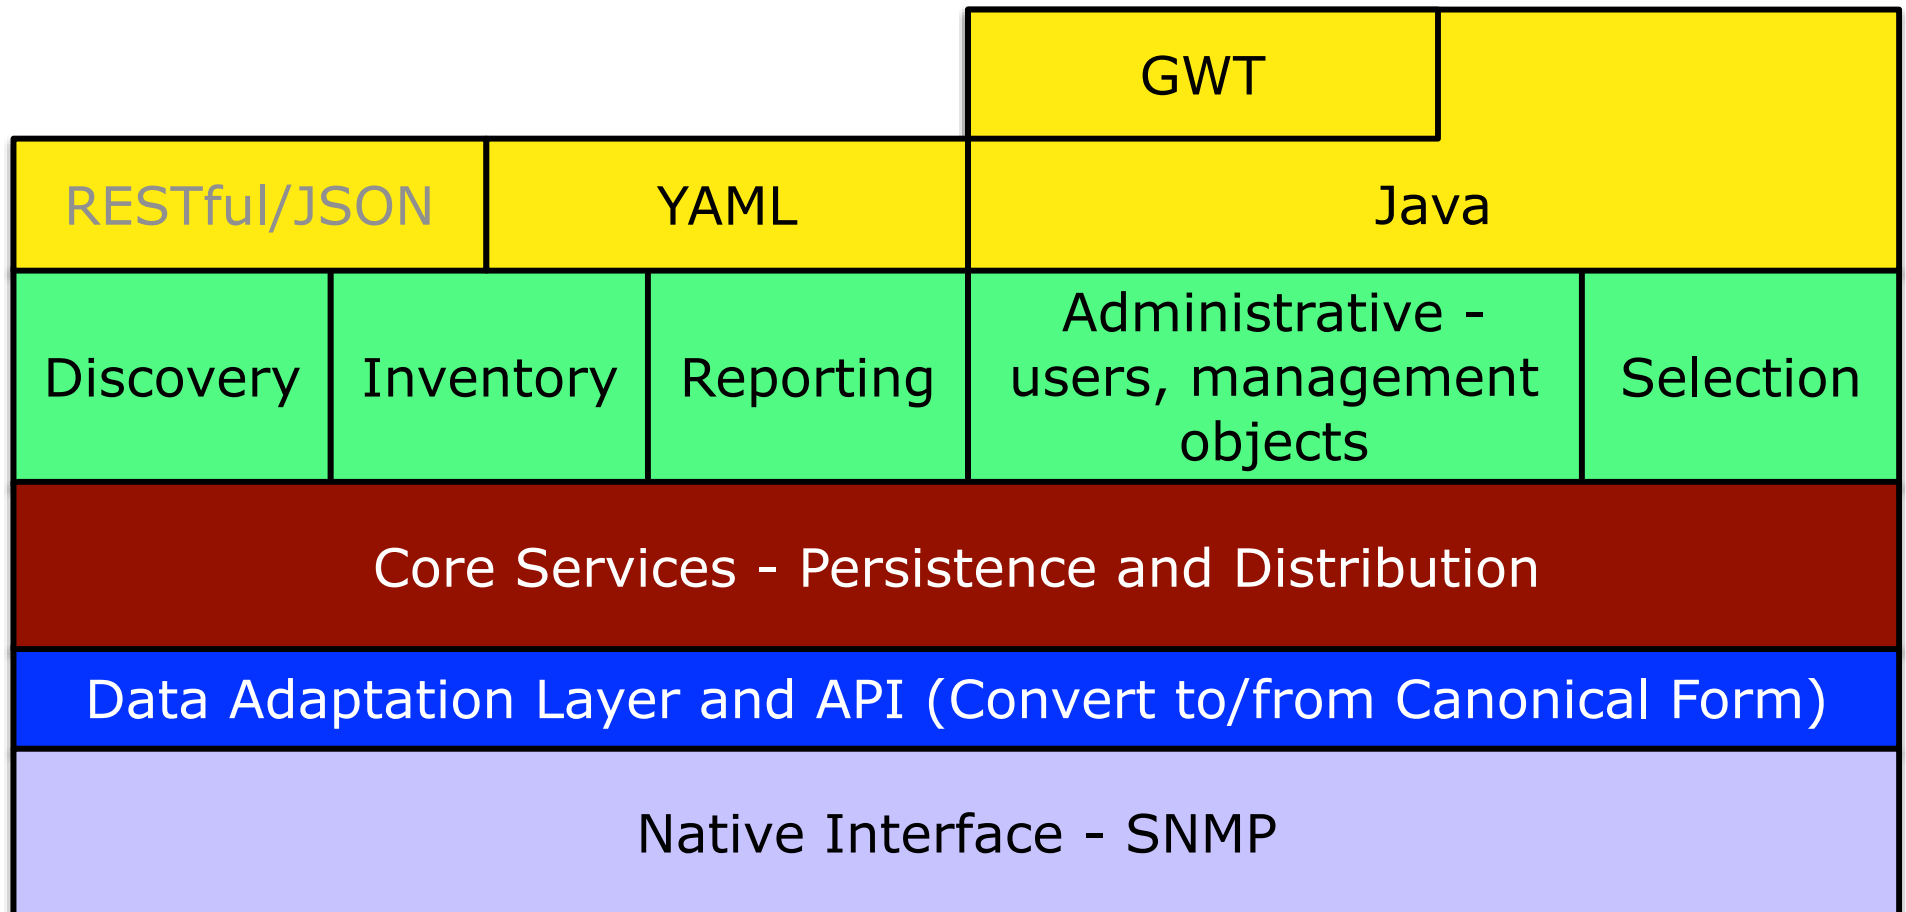
Integer Data - an API Centric View



Server Logical Design



Discovery Release - APIs, Services, Data Adaptation and Native Interfaces



APIs and Interfaces

- Programmer Centric
 - GWT
 - Java
- System administrator based:
 - RESTful/JSON
 - YAML
- Others based on user input.

Integer Services and Managers

- Services are singleton beans.
- Managers - stateless session beans.
- As many managers of a specific type as are needed/
configured, e.g.,
 - Topology.
 - Service Element Discovery.
- Aids scaling.

Core Services - Persistence

- All access to Integer data.
- Access only via service, no direct DB access.
- Access implemented with DAO's on top of Hibernate.
- Storage via MariaDB.
- MariaDB has Cassandra storage engine.
- Cassandra will store stats.
- Galera cluster for synchronous multi-master MariaDB.
 - Active-active multi-master.
 - Read and write on any node.
 - Can operate across data centers.

Core Services - Distribution

- Distribution service for systems and functions they perform.
- Keeps track of services and distributed elements.
- Each system is separately configured.
- Keeps track of state and receives messages from distributed systems.
- Uses native WildFly mechanisms for messages.

Discovery Service/Managers

- Service controls:
 - Start/stop of service element discovery.
 - Start/stop of topology discovery.
 - Per sub-net/discovery rule.
- Processing 'rules'
 - Discovery rule
 - ipTopologySeed
 - Global credentials
 - snmpLocalCredentials
 - CalendarPolicy
- State/status of running discovery.

Discovery Service

```
/**
 * The Interface DiscoveryServiceInterface.
 *
 * @author David Taylor
 */
public interface DiscoveryServiceInterface extends BaseServiceInterface {

    /**
     * Start a discovery with the given DiscoveryRule.
     *
     * @param rule
     * @return
     * @throws IntegerException
     */
    public DiscoveryId startDiscovery(DiscoveryRule rule) throws IntegerException;

    /**
     * @param discoveryId
     * @throws IntegerException
     */
    void discoveryComplete(DiscoveryId discoveryId) throws IntegerException;

    /**
     * @param id
     * @param errorCode
     * @param args
     */
    void discoveryError(DiscoveryId id, NetworkErrorCodes errorCode,
        DisplayableInterface[] args);

    /**
     * Save ServiceElement
     * @param accessElement
     */
    void discoveredServiceElement(ServiceElement accessElement);

    /**
     * Stop the running discovery specified by the DiscoveryId
     *
     * @param id
     */
    void stopDiscovery(DiscoveryId id);
}
```

API To ServiceElements

```
public interface ServiceElementAccessManagerInterface extends
    BaseManagerInterface {

    /**
     * Add or update a service element. If the service element does not exist in
     * the database. Then a new service element will be created. If the service
     * element exists in then the service element will be updated.
     *
     * @param serviceElement
     * @return Updated ServiceElement
     * @throws IntegerException
     */
    ServiceElement updateServiceElement(ServiceElement serviceElement)
        throws IntegerException;

    /**
     * Find the ServiceElements that have the given parent ID. This is the list
     * of service elements that are children of the parent service element. Ex A
     * port service element would have one or more child interface service
     * elements.
     *
     * @param parentId
     * @return
     * @throws IntegerException
     */
    ServiceElement[] getServiceElementByParentId(ID parentId)
        throws IntegerException;
```

Inventory Service

- Controlled by inventory rules:
 - What service elements are covered by the rule:
 - Location.
 - Criticality.
 - Technologies (e.g, routers, name servers, etc.).
 - Changes since last run.
 - New systems.
 - Missing systems.
 - Definition of:
 - What it means to be 'missing' - number of runs absent.
 - Number of runs a new system must exist before being added.
 - Notification actions.
 - Systems to exempt (e.g, those with ifAdmin status down).
 - What a change is for a specific type of service element (based on unique identifiers for service element types).

Reporting Service

- Same selection method as other parts of the system.
- For this release:
 - Entire inventory.
 - Changes or other selections.
- Basic format/text report.
- JSON output.
- Email.
- Event/alarm window.

Selecting Technologies

```
/**
 * Get the service identified by the given ID
 *
 * @param service
 * @return Service for the given ID
 * @throws IntegerException
 */
Service getServiceById(ID service) throws IntegerException;
```

```
/**
 * Get all the services that are in the database.
 *
 * @return List of Services found in the database.
 * @throws IntegerException
 */
Service[] getAllServices() throws IntegerException;
```

```
...
```

Selection Service

- Same approach for interacting with Integer for all functions, e.g.,:
 - What you want to see in a report.
 - What you want to see on the screen.
 - What to discover.
- Elements of a selection
 - The systems/services.
 - Should topology information be included.
 - The information about the selected systems/services of interest
 - State and (configuration, discovery, fault, etc.).
 - Utilization and other historic information.
 - Capacity.
 - Calculated values input by users.
 - How to present the information, e.g.,
 - As a topology map.
 - A tabular report.
 - Output to CSV or JSON.

Other Services/Managers

- Events/Alarms and logs.
- User and roles.
- Information for the data adaptation layer.

Data Adaptation Layer

- Maps between device and protocol details to Integer.
- Data driven.
- User modifiable.
- Consistent user experience.

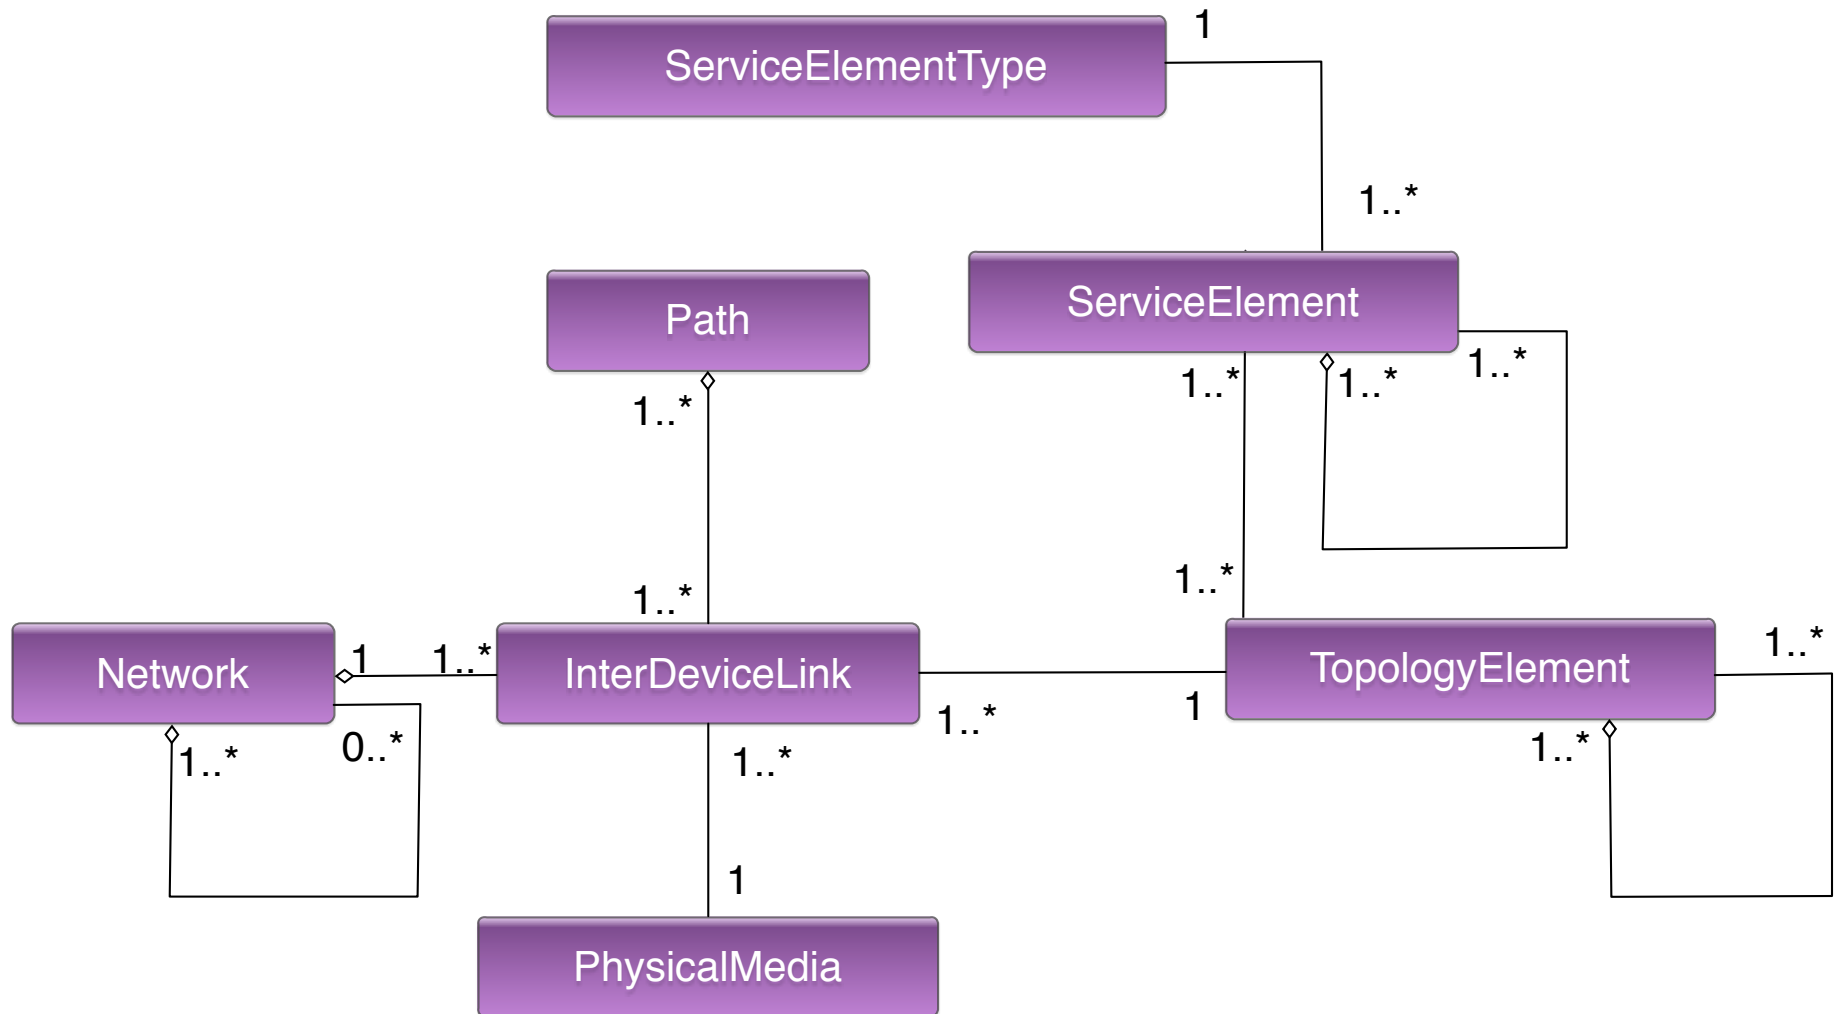
Object Hierarchies

- Main conceptual elements
 - Business services and technologies.
 - Providers and organizations.
 - Users, groups, roles.
 - Service elements and multi-layer/technology topology.
- Others used for various functions and services:
 - Discovery.
 - User selections.
 - Reporting.

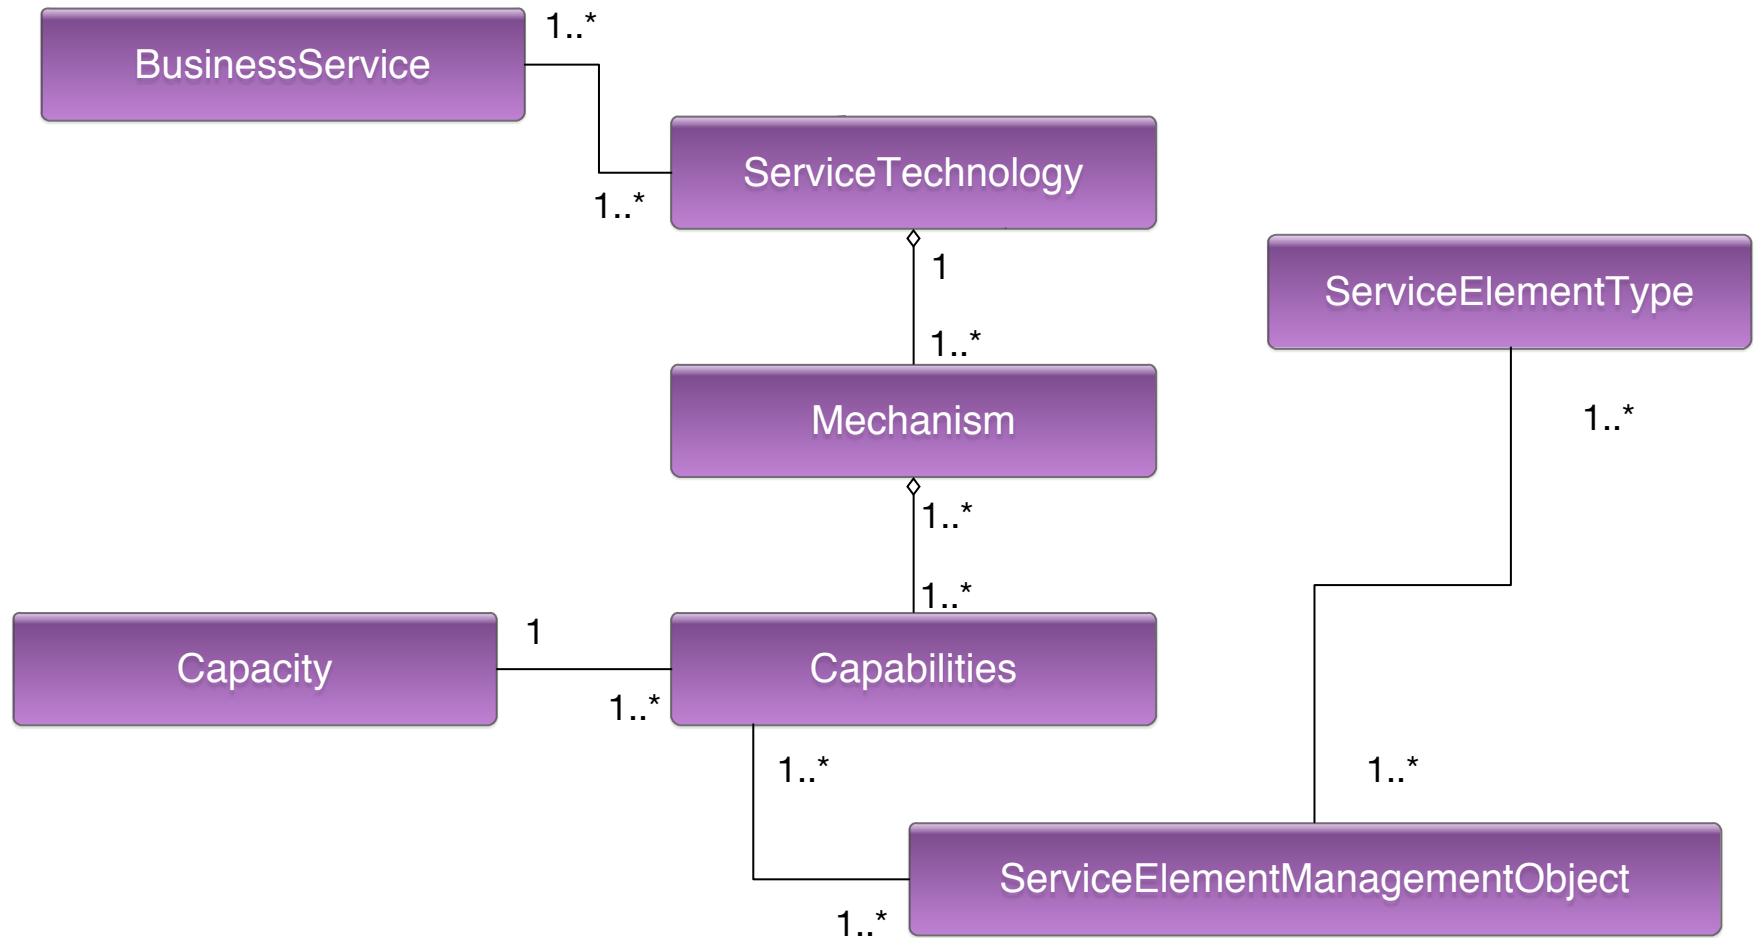
Integer Domain Knowledge

- How Integer ‘learns’ about and normalizes to a canonical form, different types of management objects:
 - SNMP - standard and vendor-specific.
 - A variety of CLIs.
 - RESTful and other interfaces and data formats used on them
 - Different APIs/interfaces for different environments like AWS.
 - DevOps and automation systems like Puppet, Chef, etc.
- How Integer ‘learns’ what makes a particular device, or part it contains unique in terms of what it can do.
- How Integer connects the large number of combinations into knobs that a human can control and understand and use to train Integer further.

Systems and Topology



Domain Knowledge Enables Service Management

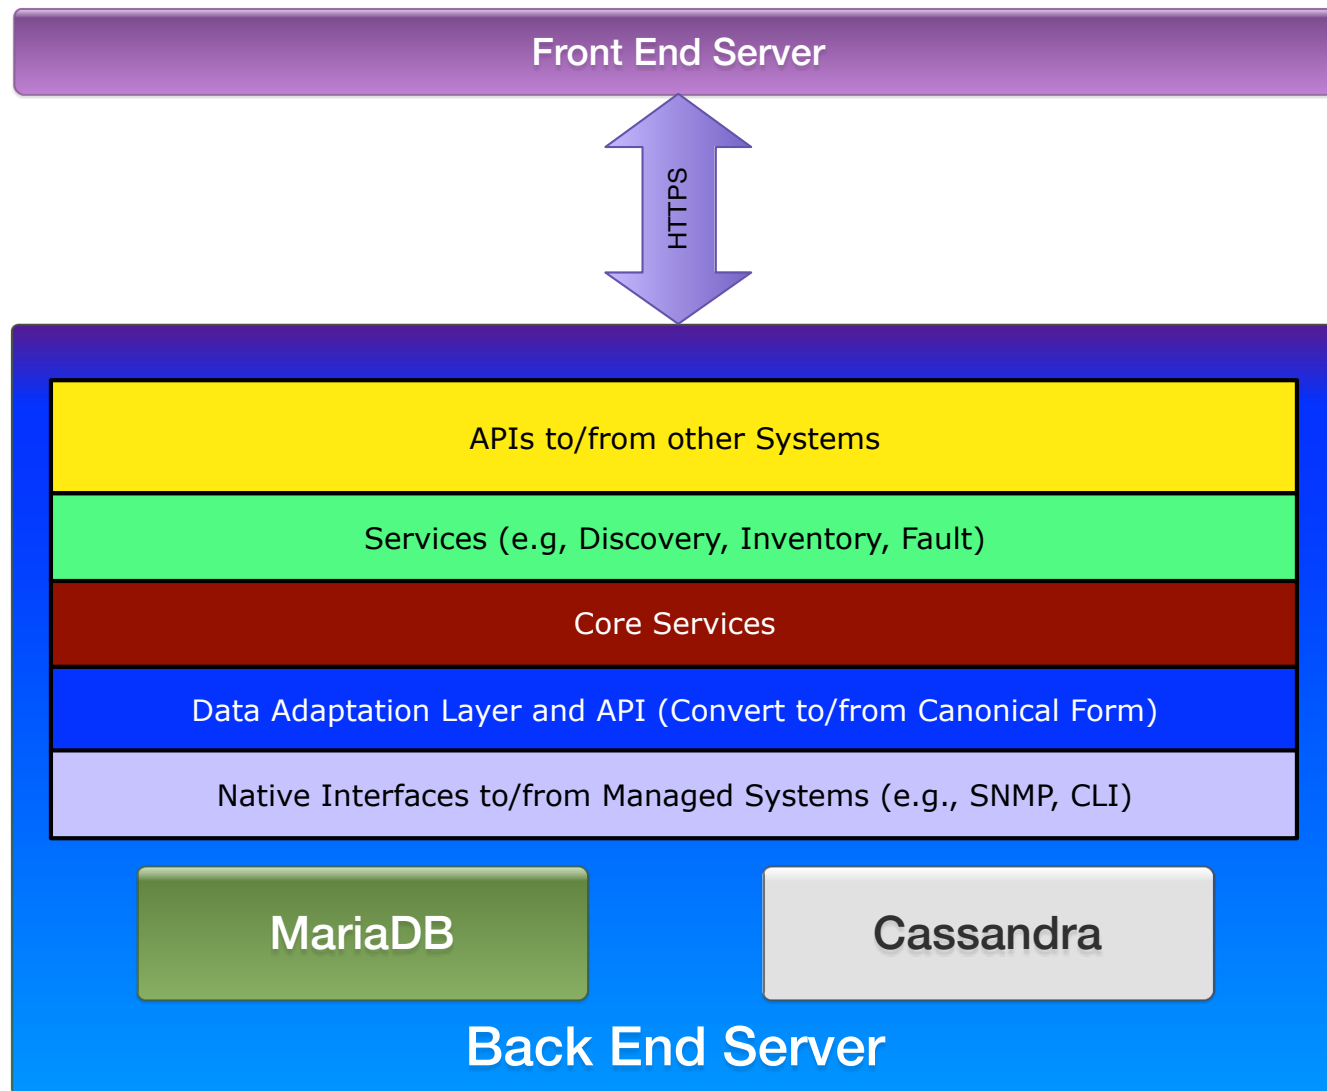


YAML Example

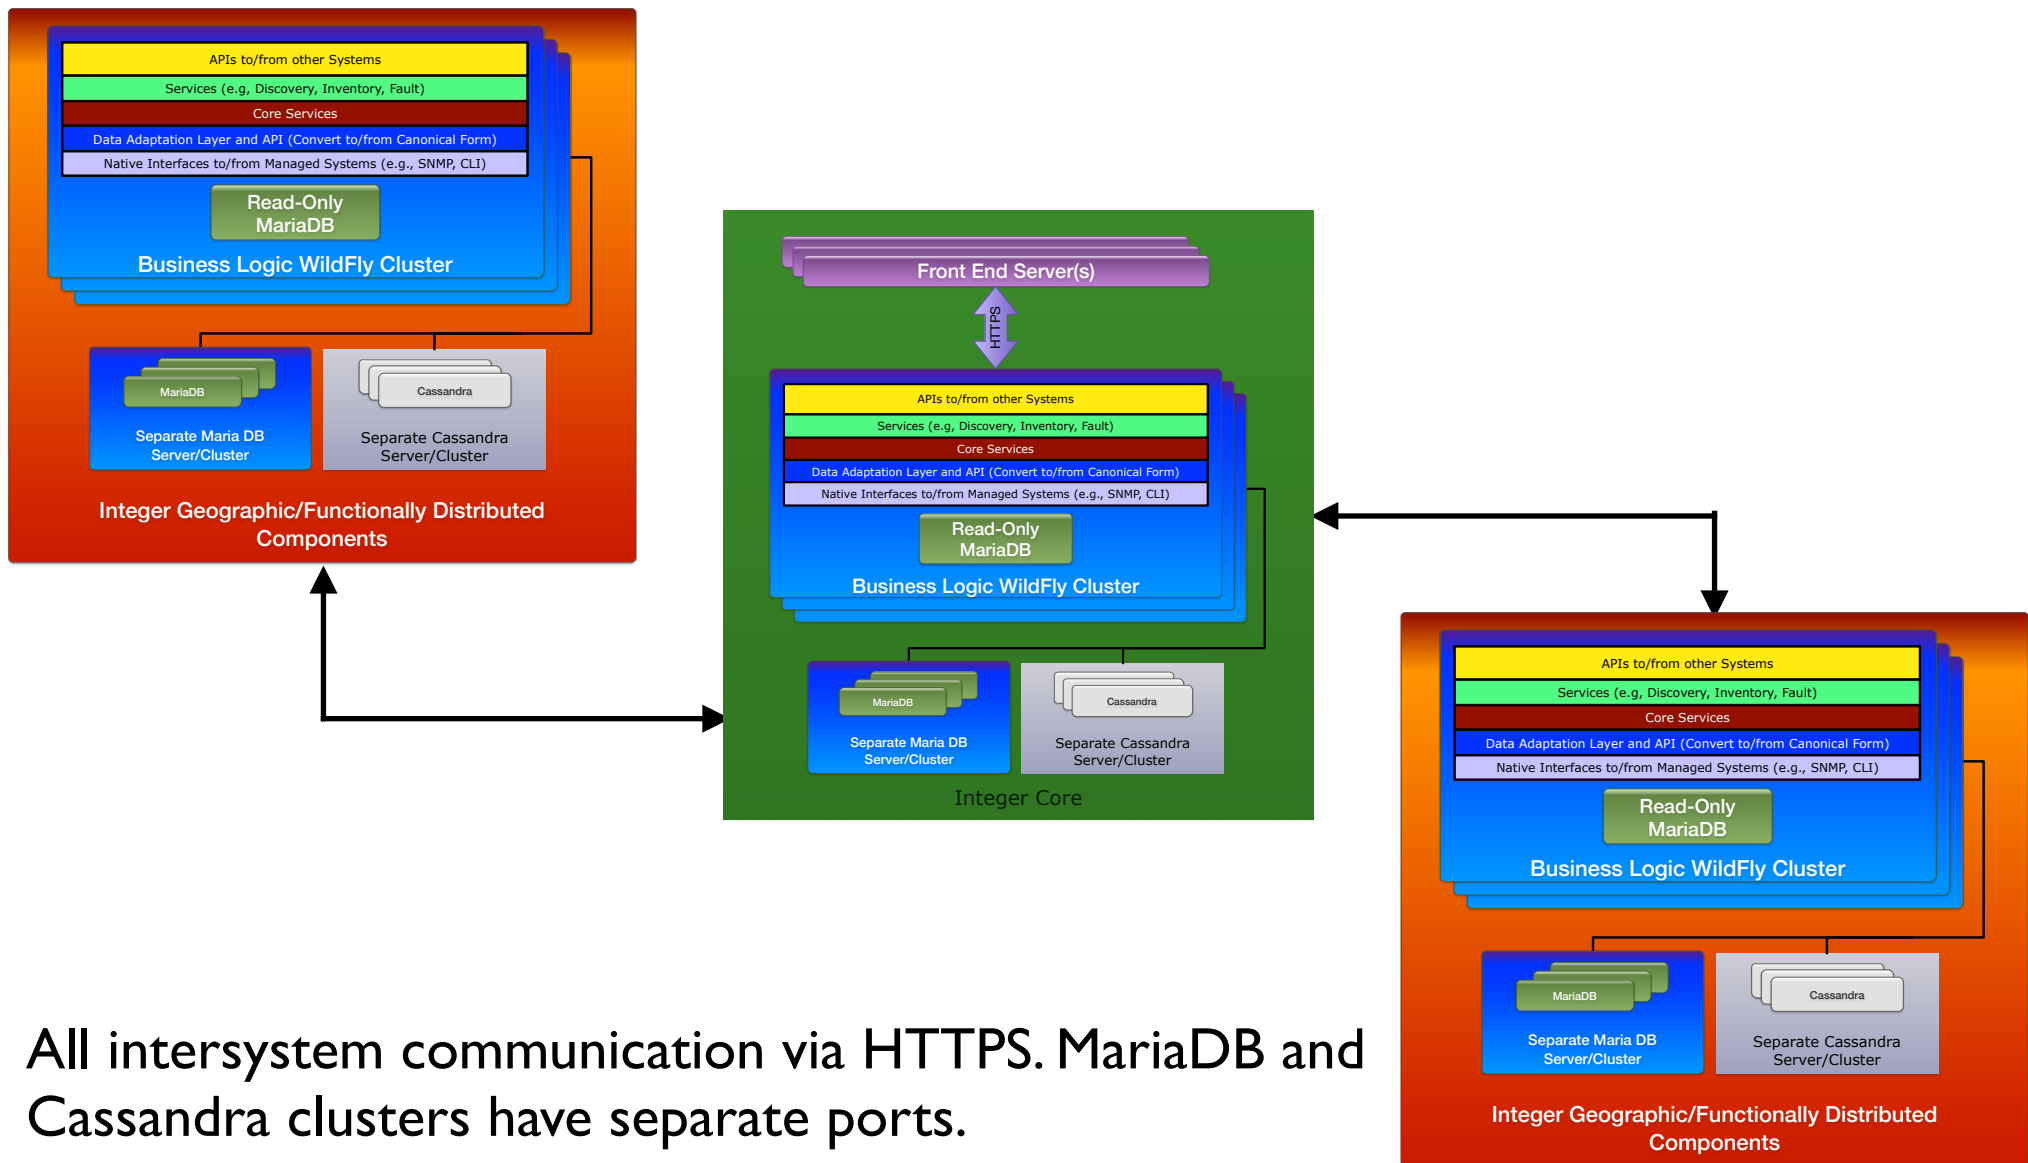
```
1  ---
2
3  # management definition for CDP
4
5  serviceElementTypes:
6
7  - name: cdp-interface
8    description: Cisco enterprise CDP MIB, per-interface attributes.
9    vendor: cisco
10   extend: interface
11   accessMethod:
12     protocol: SNMP
13     mib: CISCO-CDP-MIB
14     mapping: direct
15   serviceElementTypeTranslates:
16   - name: cpdInterface
17     mapping: direct
18     category: interface
19
20   managementObjects:
21
22   - name: cdpInterfaceEnable
23     uri: cdpInterfaceEnable
24     capability: cdp-interface-enable
25
26   - name: cdpInterfaceName
27     uri: cdpInterfaceName
28     capability: cdp-interface-name
29
30   - name: ifMtu
31     uri: ifMtu
32     extension: exclude
33
```

Deployment and Operation

Distribution - Simple Deployment



Distribution - Large/Complex Case



All intersystem communication via HTTPS. MariaDB and Cassandra clusters have separate ports.

Discovery Release Objectives

- Covers SNMProwl +.
- Complete inventory of the environment.
 - Network elements - including layer 2 and message stream modification systems
 - Servers.
 - Interconnections.
 - Details of the 'contents' of each system - e.g., cards, ports, etc.
 - Some high-level software.
- Automated system for detecting, tracking and reporting changes.
- Feed to systems like Nagios that require 'inventory' information.

Discovery Deployment

- Single system or distributed.
 - Configured - scope of subnets included in the discovery.
 - Distributed elements roll up to the main installation.
 - Each can have a separate set of 'rules' that control behavior.
- Virtualized, physical, or cloud-based environments.
- Each 'rule' may have a different schedule if desired.
- Systems may have multiple tasks discovering devices and topology.

SNMP Background

- SNMP.
- Standard.
- Widely implemented.
- Familiar.
- Strengths and weaknesses understood.

SNMP High-Level Details

- Multiple versions.
- Data stored in a Management Information Base - MIB.
- Language for defining objects:
 - Structure of Management Information (SMI).
 - “Imperfect” subset of Abstract Syntax Notation.1.
- Protocol: query functions, ‘set’ function, asynchronous messages.
- Administrative framework.

Discovery Mechanism

- First release SNMP-based.
- Find devices that will respond to SNMP requests.
- Identify vendor and other information.
- Identify device containment hierarchy.
- Retrieve interface information.
- Assign a ServiceElementType.
- Create systems and associated network information.

Discovery - Service Element Discovery

- Determine range of hosts.
- For each system found in range we attempt to retrieve 6 elements from the System Group we use to:
 - Establish vendor and top level information.
 - Determine device's containment structure.
 - Learn about each sub-element.
 - Learn about interfaces and network connections.
- Learn about other technologies.
- Save to database.
- Repeat for each network specified.

Discovery - Topology

- Use seed information from Discovery Rules.
- Information from device discovery (if run previously).
- Discovery of different technologies relevant to topology and connectivity:
 - CDP.
 - LLDP.
 - Routing information by protocol if desired:
 - Static.
 - BGP.
 - OSPF.
 - RIP.
 - Etc.
 - Other types of interconnection such as router or other types of redundancy.
 - Layer 2 (VLANs).
 - Layer 2.5 like MPLS.
- Discover additional networks based on rules.
- Save information to database.

Discovery - Next Steps

- User interface for domain knowledge.
- Manual insertion/placement of layer 1 elements.
- Extend to service dependencies.
- Extend to storage.
- Extend to cloud (e.g., AWS).
- Extend to virtualization.

integer.harvard.edu

Additional Information

- info@integer.harvard.edu
 - Specific requests for information or to make comments.
- integerinfo-subscribe@integer.harvard.edu
 - Mailing list used periodically to send out news and updates on the initiative.
- integerproject-subscribe@integer.harvard.edu
 - Subscription mailing list for people to discuss the initiative.
- Become a ‘tester’.
- Contribute other skills and knowledge.
- publictest.integer.harvard.edu
 - Coming soon.